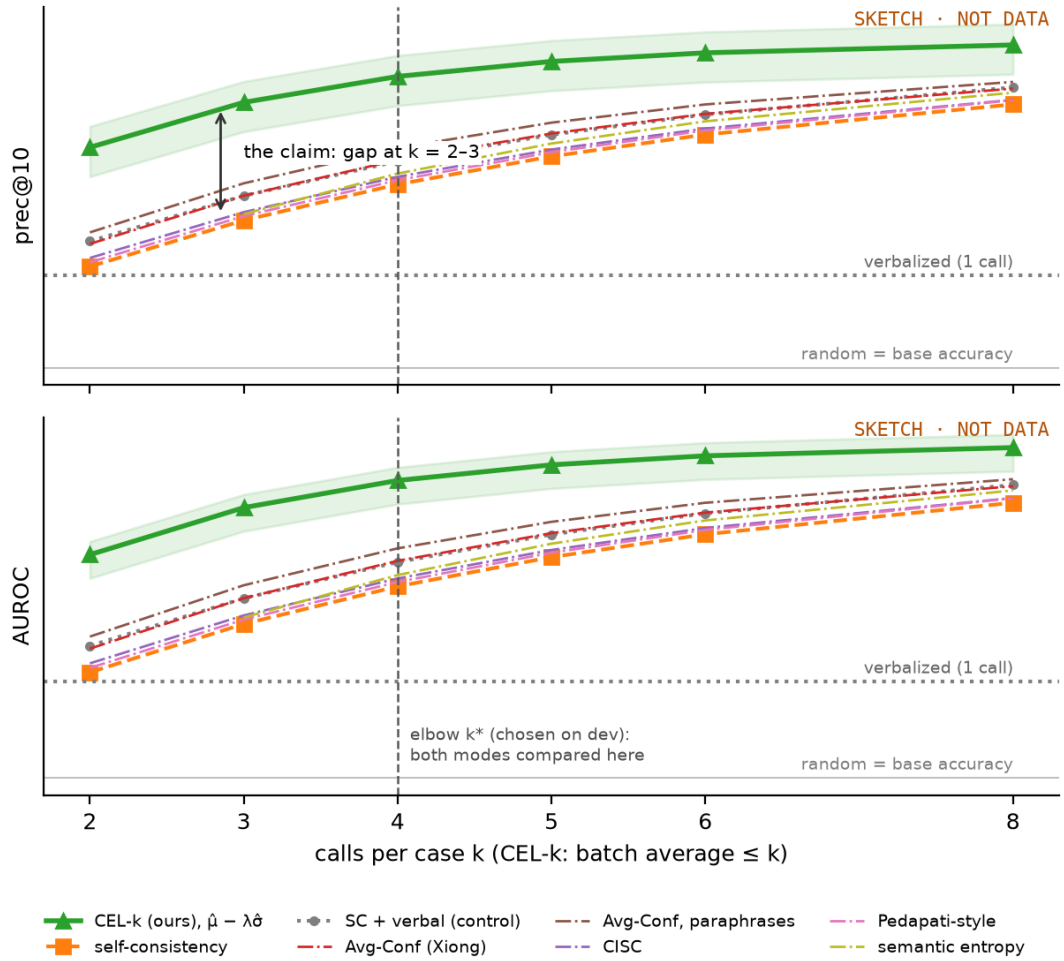
CEL: expected results

2026-10-13

No real data has been collected yet: every figure is a sketch of the pre-registered prediction (*SKETCH · NOT DATA*, no values on the quality axes). Each result is computed per (dataset, model) cell on 500 eval items, 10 batches of 50, top-10 per batch; everything is fitted on the dev slice. The two modes are defined in `method.md`.

R1. Budget curve (CEL-k, ranking mode)

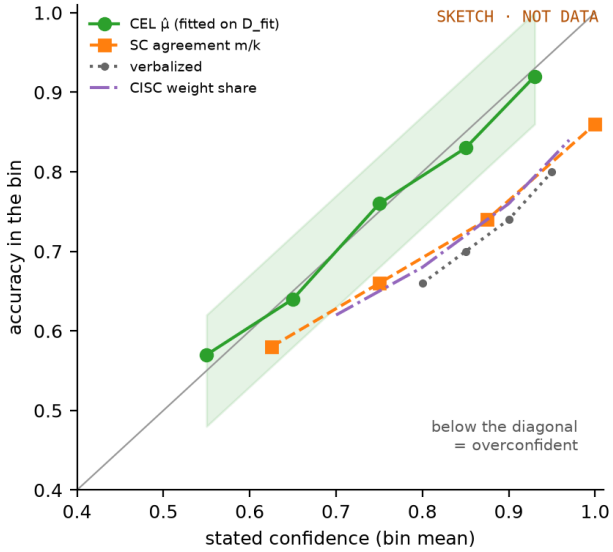


prec@10 and AUROC against calls per case. CEL-k (dev-chosen plan, ranked by $\hat{\mu} - \lambda\hat{\sigma}$) spends k on average over a batch; every baseline spends exactly k. The claim is the gap at k = 2–3, and the deciding comparison is the SC + verbal control (H3). The elbow k* is picked on the dev curve, never on eval.

R2. Metrics table at the elbow (both modes)

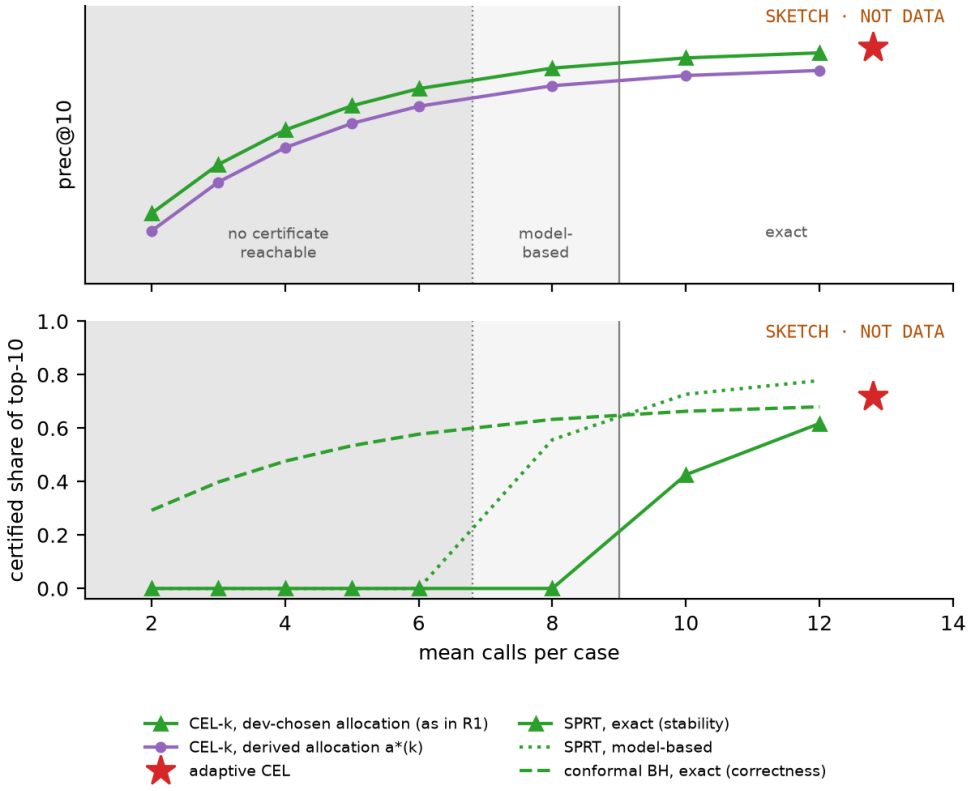
At k*, CEL-k and every baseline: prec@10, calls, accuracy, AUROC, AUARC, Δ vs CEL-k, and conformal power at FDR $\alpha = 0.1$ (`elbow.tex`). Certified CEL sits beside them at its own spend, since the SPRT sets it, with its certified share (`elbow.tex`, `main_ladder.tex`, `certified.tex`). The single confirmatory test, CEL against matched self-consistency, is in `confirmatory.tex`.

R3. Calibration (ranking mode)



Reliability of every score that claims to be P(correct), at k = 3 and at the matched budget. $\hat{\mu}$ should sit on the diagonal; verbal confidence and agreement should be overconfident, worst where stable-but-wrong answers pile up.

R4. Mode linkage



CEL-k as k grows, and certified CEL as a point. Bottom: the share of the top-10 that is certified. The SPRT certificate is zero below its budget B* and rises after it; conformal control works at every k. The CEL-k curve should meet certified CEL at its mean spend.